

Candidate Intelligence — Master Event & Signal Map

Reconciliation of old model → corrected model. Events, observations, and signals.

Diane Wilkinson | AI-Native TA Operations Portfolio | dianewilkinson.github.io

1. The Three Layers — Architecture Reminder

Every piece of data in this system lives in exactly one of three layers. Understanding which layer something belongs to is what keeps the model clean.

Layer	What it is	Example
Event	Something that happened. A moment in time. A verb.	application_submitted, interview_completed, task_submitted
Observation	Something measured at that moment. A noun with a value. No scoring.	resume_typo_count = 3, response_time_hours = 4.2, pages_viewed = 7
Signal	What the observation means, normalized 0-1 with a formula and direction.	typo_free_signal = 0.7 (derived from typo_count via formula)

2. Event Model — Old vs Corrected

Reconciliation complete as of March 2026. Old model rows marked RED are removed from the active event model. This table is the audit trail showing what changed and why.

Every event from the old data model is accounted for below. Green = kept unchanged. Yellow = renamed or restructured. Red = removed.

	Green = Kept (name unchanged)
	Yellow = Renamed or restructured
	Red = Removed from event model

2.1 Candidate-Facing Events (generate candidate signals)

Event Name	Status	Candidate Signals?	Notes
application_submitted	KEPT	Yes	Absorbs all resume and agent research observations. Opt-in flags recorded here.
resume_parsed	REMOVED	No	Folds into application_submitted. Not a separate event.
interview_invited	NEW	Yes	Timestamp only. Starts responsiveness clock.

interview_rsvp	NEW	Yes	Replaces email_received. response_time_hours = delta from interview_invited timestamp.
email_sent	REMOVED	No	Too generic. Replaced by named candidate events.
email_received	REMOVED	No	Too generic. Replaced by interview_rsvp.
follow_up_sent	REMOVED	No	Replaced by interview_thank_you_sent and candidate_outreach.
interview_thank_you_sent	NEW	Yes	Candidate opt-in post-interview follow-up. Boolean + message quality observations.
candidate_outreach	NEW	Yes	Unprompted candidate-initiated contact. is_unprompted always True by definition - field removed as redundant.
portal_session	RENAMED	Yes	Replaces portal_login. Covers authenticated and unauthenticated - same signal weight either way.
portal_login	REMOVED	No	Renamed to portal_session.
site_engage	REMOVED	No	Collapsed into portal_session - authenticated and unauthenticated carry same signal weight.
task_assigned	KEPT	No	Timestamp only. Required to calculate days-to-submission timeliness signal.
task_completed	REMOVED	No	Renamed to task_submitted.
task_submitted	RENAMED	Yes	Replaces task_completed for required tasks. Carries task_quality_raw, completed_on_time, response_time_hours, typo_count, grammar_errors.
opt_in_submission	NEW	Yes	ALL opt-in submissions. submission_type field: video_intro, ai_interview, portfolio, assessment, task, research. Choice = Motivation; quality = Capability; care = Conscientiousness.
live_prep_requested	NEW	Yes	Candidate requests live prep with recruiter. Opt-in. Motivation + Conscientiousness.
live_prep_completed	NEW	Yes + Workflow	Confirms prep happened. Candidate engagement signal + workflow timestamp.
interview_completed	KEPT	Yes	verbal_communication_raw from AI notetaker. interviewer_rating_raw excluded from construct scores - Decision Quality only. no_show boolean.
survey_submitted	NEW	No	Candidate experience + Workflow Intelligence only. No construct scoring.
offer_sent	KEPT	No	Workflow timestamp only.
offer_accepted	KEPT	No	Workflow timestamp only.
offer_declined	KEPT	No	Workflow timestamp only.

candidate_rejected	KEPT	No	rejection_stage recorded. Decision Quality + Workflow Intelligence.
--------------------	------	----	---

2.2 Workflow-Only Events (no candidate signals)

Event Name	Status	Candidate Signals?	Notes
interview_scheduled	KEPT	No	First scheduling only. Workflow timestamp. Time in stage feeds Workflow Intelligence.
interview_rescheduled	NEW	No	Separate event per reschedule. Each gets own timestamp and running count. Cleaner than reschedule_count on a single event.
recruiter_screen	KEPT	No	Stage timestamp. Workflow only.
hm_interview	KEPT	No	Stage timestamp. Workflow only.
final_interview	KEPT	No	Stage timestamp. Workflow only.
hiring_committee	KEPT	No	Internal deliberation. Workflow only.
recruiter_interview_prep	NEW	No	Recruiter conducts prep session. Workflow timestamp. Candidate opt-in captured on live_prep_requested.
survey_requested	NEW	No	System sends survey. Workflow trigger. Response captured on survey_submitted.

3. Observation Fields — Old vs Corrected

Every observation field from the old model is accounted for below. Fields marked REMOVED are either deprecated, replaced, or never should have been observations (e.g. keyword matching).

3.1 Application Stage Observations (at application_submitted)

Observation Field	Event	Layer	Notes / Disposition
resume_typo_count	application_submitted	KEPT	Raw count. Feeds typo_free_signal.
resume_grammar_error_count	application_submitted	KEPT	Raw count. Feeds grammar_quality_signal.
resume_syntax_issue_count	application_submitted	KEPT	Raw count. Feeds syntax_quality_signal.
resume_word_count	application_submitted	KEPT	Feeds resume_clarity_signal via optimal range formula.
resume_formatting_raw	application_submitted	KEPT	Raw formatting score. Feeds resume_formatting_signal.
resume_keyword_match_ratio	application_submitted	REMOVED	All resumes have degraded keyword signal. Not scored. Removed from model.

time_on_application	application_submitted	NEW	Minutes spent completing the application. Feeds Motivation.
opt_in_video_intro	application_submitted	NEW	Boolean. Did candidate submit optional video intro.
opt_in_ai_interview	application_submitted	NEW	Boolean. Did candidate submit optional AI interview.
opt_in_assessment	application_submitted	NEW	Boolean. Did candidate submit optional assessment at application stage.

3.2 Screening Agent Research Observations (at application_submitted, agent-generated)

Observation Field	Event	Layer	Notes / Disposition
years_experience_actual	application_submitted	KEPT	Agent-verified, not self-reported. From employer research.
years_experience_required	application_submitted	KEPT	From JD / rubric configuration.
tech_stack_skills_required	application_submitted	KEPT	From JD / rubric. Number of required skills.
tech_stack_skills_matched	application_submitted	KEPT	Agent-verified match count.
tech_stack_avg_exposure_years	application_submitted	KEPT	Agent-verified depth per matched skill.
employer_industry_match	application_submitted	KEPT	Boolean. Agent-researched.
employer_size_match	application_submitted	KEPT	Boolean. Agent-researched.
role_title_similarity_cosine	application_submitted	KEPT	Agent derives from JD research at prior employers, not title string matching.
public_evidence_found	application_submitted	NEW	Boolean. Agent found public portfolio, articles, speaking engagements.
public_evidence_type	application_submitted	NEW	Categorical: portfolio / article / speaking / forum_mention.
tech_stack_query_flag	application_submitted	NEW	Boolean. Unverifiable tech claim — triggers clarification email, does not score.
linkedin_rec_trading_detected	application_submitted	NEW	Boolean. If True, recommendation signal silently nulled. Not surfaced.

3.3 Responsiveness Observations

Observation Field	Event	Layer	Notes / Disposition
-------------------	-------	-------	---------------------

response_time_hours	interview_rsvp	KEPT	Delta between interview_invited and interview_rsvp timestamps. Replaces email-based measurement.
response_time_hours (old email_sent)	email_sent	REMOVED	Email events removed. Responsiveness now measured via interview_invited → interview_rsvp.
response_time_hours (old email_received)	email_received	REMOVED	Email events removed.
response_time_hours (old recruiter_screen)	recruiter_screen	REMOVED	Workflow event — response time here was conflating process efficiency with candidate behavior.

3.4 Portal & Site Engagement Observations

Observation Field	Event	Layer	Notes / Disposition
session_duration_minutes	portal_session	KEPT	Renamed event (was portal_login). Same field.
pages_viewed	portal_session	KEPT	Renamed event. Same field.
content_downloaded_count	portal_session	KEPT	Renamed event. Same field.
prep_material_accessed	portal_session	KEPT	Boolean. Renamed event. Same field.
time_on_site_minutes	site_engage	NEW	Unauthenticated visit duration.
pages_viewed_unauthenticated	site_engage	NEW	Softer signal — we know less about who this is.

3.5 Task Observations

Observation Field	Event	Layer	Notes / Disposition
task_quality_raw	task_submitted	KEPT	Renamed event (was task_completed). Same field.
completed_on_time	task_submitted	KEPT	Renamed event. Same field.
response_time_hours	task_submitted	NEW	Days between task_assigned and task_submitted. Timeliness signal.
task_quality_raw	opt_in_task_submitted	KEPT	Same field, different event. Weighted differently — opt-in carries Motivation signal too.
submission_type	opt_in_submission	NEW	Categorical: video_intro / ai_interview / portfolio / assessment / research.
quality_raw	opt_in_submission	NEW	Quality score for the submission. Feeds Capability.

3.6 Communication Observations

Observation Field	Event	Layer	Notes / Disposition
message_sentiment_raw	interview_thank_you_sent	KEPT	Sentiment of post-interview follow-up message.
typo_count	interview_thank_you_sent	NEW	Attention to detail in follow-up message.
grammar_error_count	interview_thank_you_sent	NEW	Written quality in follow-up.
response_time_hours	interview_thank_you_sent	NEW	How quickly after interview was thank you sent.
message_sentiment_raw	candidate_outreach	KEPT	Sentiment of unprompted outreach.
typo_count	candidate_outreach	NEW	Attention to detail in outreach message.
is_unprompted	candidate_outreach	REMOVED AS FIELD	Always True by definition of this event — field is redundant. Event name carries the signal.
is_follow_up	follow_up_sent	REMOVED	Event removed. Replaced by named events above.
is_unprompted	follow_up_sent	REMOVED	Event removed.
message_sentiment_raw (follow_up_sent)	follow_up_sent	REMOVED	Event removed.

3.7 Interview Observations

Observation Field	Event	Layer	Notes / Disposition
interviewer_rating_raw	interview_completed	KEPT	0-5 scale. Feeds interviewer_impression_signal. Excluded from construct scores — Decision Quality only.
no_show	interview_completed	KEPT	Boolean. If True, no_show_free_signal = 0.
reschedule_count	interview_scheduled	KEPT	Workflow event. Feeds scheduling_reliability_signal but note: this is workflow-side data not candidate behavioral event.

3.8 Fields Kept as Context / Metadata (not scored)

Observation Field	Event	Layer	Notes / Disposition
source_channel	application_submitted	CONTEXT	Where candidate came from. Feeds Workflow Intelligence source quality analysis. Not a candidate score input.

department	application_submitted	CONTEXT	Used for department percentile calculation and cohort grouping.
role_level	application_submitted	CONTEXT	Used for cohort grouping and rubric configuration.
applied_date	application_submitted	CONTEXT	Used for time-in-stage calculations.
rejection_stage	candidate_rejected	CONTEXT	Where in the process rejection occurred. Decision Quality and Workflow Intelligence.

4. Signal Reconciliation — Old vs Corrected

Every signal from the old model is accounted for. Signals marked REMOVED either derived from removed observations or were never valid given the new philosophy.

4.1 Signals Kept — Capability Construct (40%)

Signal	From Observation	Construct	Notes
experience_fit_signal	years_experience_actual / required	Capability	Kept. Sourced from agent research not self-report.
tech_stack_overlap_signal	tech_stack_skills_matched / required	Capability	Kept. Agent-verified.
tech_depth_signal	tech_stack_avg_exposure_years	Capability	Kept. Agent-verified.
employer_context_signal	employer_industry_match + employer_size_match	Capability	Kept.
role_similarity_signal	role_title_similarity_cosine	Capability	Kept. Now derived from JD research not title string match.
resume_relevance_signal	resume_keyword_match_ratio	REMOVED	Keyword matching removed from model. AI resumes degraded this signal entirely.
task_quality_signal	task_quality_raw	Capability	Kept. Applies to both task_submitted and opt_in_task_submitted.

4.2 Signals Kept — Conscientiousness Construct (25%)

Signal	From Observation	Construct	Notes
typo_free_signal	resume_typo_count + typo_count across events	Conscientiousness	Kept. Now aggregated across all written touchpoints not just resume.
grammar_quality_signal	grammar_error_count across events	Conscientiousness	Kept. Same expansion.

syntax_quality_signal	resume_syntax_issue_count	Conscientiousness	Kept.
timeliness_signal	completed_on_time	Conscientiousness	Kept.
scheduling_reliability_signal	reschedule_count	Conscientiousness	Kept. Note: sourced from workflow event — flag for review.
no_show_free_signal	no_show	Conscientiousness	Kept.
follow_through_signal	composite	Conscientiousness	Kept. Composite of timeliness, scheduling reliability, no_show_free.
response_delay_signal	response_time_hours	Conscientiousness	Kept. Now sourced from interview_rsvp not email events.
prep_investment_signal	prep_material_accessed	Conscientiousness	Kept.

4.3 Signals Kept — Communication Construct (20%)

Signal	From Observation	Construct	Notes
responsiveness_signal	response_time_hours at interview_rsvp	Communication	Kept. Source corrected — now from interview_rsvp not email_received.
message_sentiment_signal	message_sentiment_raw across events	Communication	Kept. Aggregated across interview_thank_you_sent and candidate_outreach.
resume_clarity_signal	resume_word_count	Communication	Kept.
resume_formatting_signal	resume_formatting_raw	Communication	Kept.
interviewer_impression_signal	interviewer_rating_raw	EXCLUDED	Intentionally excluded from all construct scores. Decision Quality dashboard only.

4.4 Signals Kept — Motivation Construct (15%)

Signal	From Observation	Construct	Notes
portal_engagement_signal	session_duration + pages_viewed	Motivation	Kept. Source renamed to portal_session.
session_depth_signal	session_duration_minutes	Motivation	Kept.
content_exploration_signal	pages_viewed	Motivation	Kept.
content_download_signal	content_downloaded_count	Motivation	Kept.
follow_up_signal	interview_thank_you_sent	Motivation	Kept. Source event renamed.

proactive_follow_up_signal	candidate_outreach	Motivation	Kept. Source event renamed. Always unprompted by definition.
scheduling_distance_signal	interview_scheduled timestamp delta	Motivation	Kept. Workflow event source — flag for review same as scheduling_reliability.

4.5 Signals Removed

Signal	From Observation	Construct	Notes
resume_relevance_signal	resume_keyword_match_ratio	REMOVED	Keyword matching philosophy removed entirely. AI resumes have degraded this signal.
primary_signal_value	routing field	REMOVED	Was a routing field not a true signal. Removed from scored model.
signal_domain	routing field	REMOVED	Categorical routing field. Not a signal. Removed from scored model.

5. Open Items — To Resolve

The following items were flagged during mapping and need a decision before the data model is regenerated.

Item	Question / Decision Needed
reschedule_count sourcing	Currently on interview_scheduled (workflow event). Should this move to a candidate-facing event, or is workflow-side sourcing acceptable given it still reflects candidate behavior?
scheduling_distance_signal sourcing	Currently derived from interview_scheduled timestamp delta. Same question as above — workflow event sourcing a candidate signal.
oral communication quality	Where does verbal quality from interviews get observed? AI notetaker transcript → Apps Script sentiment + quality analysis. What event does this attach to and what are the observation field names?
education & certifications	Not yet mapped to an event or observation field. Where does this get recorded? At application_submitted by the agent? Or candidate-submitted field?
references & work samples	Not yet mapped. Same question — agent-found or candidate-submitted? What event?
survey content scoring	survey_submitted event added. What observations does it carry and do any feed candidate construct scores or is it purely candidate experience / employer brand data?
site_engage authenticity	Unauthenticated site visits are softer signal. How much weight should these carry vs authenticated portal_session signals?

6. New Events Added — Session 2 (Stage 2 & 3 Mapping)

The following events were identified during stage-by-stage mapping and are additions to the model not captured in the initial reconciliation.

6.1 Candidate-Facing Events Added

Event Name	Status	Signals?	Notes
confirmation_sent	NEW (Workflow)	Candidate response	System sends interview confirmation. Candidate calendar_confirmed boolean is the response observation.
reminder_sent	NEW (Workflow)	Candidate engagement	24hr reminder before interview. Candidate open/engage feeds Motivation.
opt_in_assessment_nudge	NEW (Workflow)	Candidate response	System nudges candidate about available optional assessment. Response or submission feeds Motivation + Capability.
interview_followup_sent	NEW (Workflow)	Candidate response	Automated 1hr post-interview thank you + next steps. Sets timeline expectation. Candidate response feeds Conscientiousness.
recruiter_touchpoint	NEW (Workflow)	Candidate response	Personal recruiter update — never no-reply. Fires before expectation deadline lapses and every 7 days post-interview until decision. Candidate response feeds Conscientiousness.
candidate_questions_asked	NEW (Candidate)	Yes	Questions candidate asks during interview. Scored for question quality (Communication) and depth of preparation (Conscientiousness). AI notetaker captures.

6.2 New Observation Fields — Scheduling Stage

Observation Field	Event	Layer	Notes
calendar_confirmed	interview_scheduled	NEW	Boolean. Did candidate add interview to calendar. Feeds Conscientiousness.
scheduling_method	interview_scheduled	NEW	Categorical: instant_book vs availability_submitted. Instant book = higher Motivation signal.
scheduling_distance_hours	interview_scheduled	NEW	Hours between interview_rsvp and interview_scheduled. Large gap = lower urgency = lower Motivation.
reschedule_initiated_by	interview_rescheduled	NEW	Categorical: candidate vs recruiter vs system. CRITICAL — candidate only penalized for candidate-initiated reschedules. Recruiter or system-initiated reschedules must not affect candidate scoring.
reschedule_count_candidate	interview_rescheduled	NEW	Running count of candidate-initiated reschedules only. Derived from

			reschedule_initiated_by = candidate.
--	--	--	--------------------------------------

6.3 New Observation Fields — Interview Stage (from AI Notetaker + Apps Script)

Observation Field	Event	Layer	Notes
verbal_communication_raw	interview_completed	NEW	AI notetaker transcript analyzed by Apps Script. Applies across recruiter_screen, hm_interview, final_interview. Feeds verbal_communication_signal -> Communication construct.
question_quality_raw	interview_completed	NEW	Quality and depth of questions candidate asked during interview. AI notetaker captured. Feeds Communication + Conscientiousness.
teaching_depth_raw	interview_completed	NEW — OPEN	How far candidate went beyond surface-level answers when given teaching-mode questions. Source TBD: AI notetaker scored or recruiter-rated field.
sentiment_trajectory	interview_completed	NEW	Did candidate sentiment improve during interview? Positive trajectory = recruiter connection techniques working. Feeds candidate Communication signal.
interview_duration_minutes	interview_completed	NEW — OPEN	Length of interview. Long duration may indicate strong engagement. Construct assignment TBD: Motivation, Communication, or both.

6.4 New Signals Added

Signal	From Observation	Construct	Notes
verbal_communication_signal	verbal_communication_raw	Communication	New. Normalized 0-1 from AI notetaker analysis. Applies at every interview stage.
question_quality_signal	question_quality_raw	Communication + Conscientiousness	New. Candidate questions during interview. Depth and relevance scored.
calendar_confirmed_signal	calendar_confirmed	Conscientiousness	New. Boolean — did candidate confirm calendar invite.
opt_in_engagement_signal	opt-in submission events	Motivation	New. Composite of all opt-in behaviors — boolean per type, not additive volume.

scheduling_urgency_signal	scheduling_distance_hours	Motivation	Replaces scheduling_distance_signal. Renamed for clarity — lower distance = higher urgency = higher score.
---------------------------	---------------------------	------------	--

7. Candidate Journey Playbook — Event & Signal Alignment

The candidate journey playbook defines outbound communications and automated workflow sequences at each stage. This section maps the playbook to the event/signal model so every touchpoint is captured and candidate responses are scored.

7.1 Outbound Communication Sequence & Candidate Signal Triggers

Stage	Outbound Communication	Candidate Signal Generated	Notes
Sourcing	InMail outreach x3 (LinkedIn)	Response → interview_rsvp; no response → archived after 3rd touch	3-touch sequence. Accepted/interested = instant book link. Future interest = nurture tag + LinkedIn connect.
Application	Confirmation email (immediate)	Portal engagement → portal_session	Links to blog, careers page, Raise the Bar video. Sets expectations.
Application	Opt-in assessment nudge	Submission → opt_in_submission	If candidate hasn't opted in. Soft nudge, not required.
Recruiter Screen	Confirmation (immediate @ booking)	calendar_confirmed boolean	Sets 24hr reminder + 1hr follow-up automations.
Recruiter Screen	Reminder (24hrs before)	Open/engage → Motivation	Includes Google Meet link + rescheduling link.
Recruiter Screen	Follow-up (1hr after screen)	Candidate reply → Conscientiousness	Thank you + next steps + consent request. Sets 48hr expectation.
HM Interview	Invite (within 48hrs of screen)	interview_rsvp response time	Starts responsiveness clock. Includes HM name, title.
HM Interview	Confirmation (immediate @ booking)	calendar_confirmed boolean	Triggers confirmation + reminder + 1hr follow-up automations.
HM Interview	Prep packet (24hrs before)	prep_material_accessed → portal_session	Interview prep attachment, careers one-pager, interview process doc, HM LinkedIn link.
HM Interview	Reminder (24hrs before)	Open/engage → Motivation	Includes Google Meet + rescheduling link.
HM Interview	Follow-up (1hr after HMI)	Candidate reply → Conscientiousness	Thank you + action required: share thoughts. Optional: email HM directly.
Final Interview	Invite (2-day follow-up if not scheduled)	interview_rsvp response time	Congrats framing. Includes panel details, interviewer names and emails.
Final Interview	Confirmation (immediate @ booking)	calendar_confirmed boolean	Full panel details: interviewers, Google Meet, task instructions, prep materials.

Final Interview	Reminder (24hrs before)	Open/engage → Motivation	Simple. Google Meet link + recruiter contact for reschedule.
Any stage post-interview	Recruiter personal touchpoint (before expectation deadline)	Candidate reply → Conscientiousness	Never no-reply. Advance OR I haven't forgotten you. Every 7 days until decision.
Week 3+ no decision	Keep-warm with bottleneck owner contact info	Candidate engagement → Motivation	Accountability design — surfaces who is causing the delay.
Rejection	Rejection email (minimum 7-day delay)	None — process terminus	Dignified. GDPR consent request for EU candidates for future process consideration.
Offer	Offer letter	Candidate response time → Conscientiousness	Monthly salary, vacation days, contract type, start date, notice period, benefits.
Onboarding	Welcome email	None — process terminus for scoring	Handoff to hiring manager for 30/60/90 plan and team intro.

The recruiter never has to remember to follow up. The system handles the sequencing, the timing, and the triggers. The recruiter is present for the human moments — the system handles everything else. Every candidate always knows where they stand, even when the answer is still pending.

7.2 Swimlane Responsibilities (from Candidate Journey Playbook)

The automated operating model defines four lanes of responsibility across each stage. The ATS system owns the communication sequencing and timing. The recruiter owns the decisions and human touchpoints. The hiring manager owns the evaluation and scorecard. The candidate owns their own responsiveness and engagement.

Role	Sourcing	Screen	HM/Final Interview	Offer + Onboarding
Recruiter	3-touch outreach, HM intake, referral replies	Conduct screen, scorecard, advance/reject ATS	HM feedback, advance/reject ATS, coordinate panel	Coordinate offer, keep warm, rejection calls, ATS update
Hiring Manager	Define must-haves, share referrals	(not in screen)	Conduct interview, scorecard <24h, recruiter feedback	Close candidate call, team intro plan, 30/60/90
Candidate	(passive)	Book/reschedule, attend screen	Prep, attend interview, submit availability	Review and respond to offer, complete docs
ATS System	Live postings, source capture	Booking requests, confirmations, reminders, 2hr	Booking, confirmations, reminders, 2hr follow up, delay notices,	Approvals, offer template, e-sign, background check, trigger onboarding

		follow up, delay notices	scorecard nudges	
--	--	--------------------------	------------------	--

8. Open Items — Updated

Items resolved since initial version are marked. Remaining open items require a decision before data regeneration is complete.

Item	Status	Resolution / Question
reschedule_count sourcing	RESOLVE D	interview_rescheduled is now its own event. reschedule_initiated_by field added. Candidate only penalized for candidate-initiated reschedules.
scheduling_distance_signal sourcing	RESOLVE D	Renamed scheduling_urgency_signal. Derived from scheduling_distance_hours observation at interview_scheduled. Directional: lower distance = higher score.
oral communication quality	RESOLVE D	verbal_communication_raw observation at all interview events. Captured by AI notetaker, analyzed by Apps Script. Feeds verbal_communication_signal into Communication construct.
education & certifications	RESOLVE D	Agent-found or candidate-submitted at application_submitted. Part of Screening A[i]gent research layer.
references & work samples	RESOLVE D	Stage-gated. Public references and samples: application_submitted stage. Non-public references: final interview stage only.
survey content scoring	RESOLVE D	survey_submitted feeds candidate experience and Workflow Intelligence only. No construct scoring.
site_engage vs portal_session	RESOLVE D	Collapsed into portal_session. Authenticated and unauthenticated carry same signal weight.
teaching_depth_raw source	RESOLVE D	Renamed knowledge_depth_raw. AI notetaker scores transcript for confidence signals vs knowledge gaps relative to role. Feeds Capability (knowledge depth) and Communication (clarity, fluency). Not a recruiter-rated field.
interview_duration_minutes scoring	RESOLVE D	Scores efficiency and collaboration. Started on time = Conscientiousness. Ended on time = Conscientiousness. Concise answers = Communication. Mutual agreement to extend = Motivation + Communication. Too short or too long without agreement = negative signal.
Interview A[i]gent location	RESOLVE D	Browser extension — not embedded in ATS. Runs autonomously during interview: captures via notetaker, scores against rubric in real time, surfaces follow-up question suggestions to recruiter, writes scorecard to ATS automatically post-interview. Recruiter focuses on candidate — agent handles documentation.
Candidate questions scoring weight	RESOLVE D	question_quality_raw: thoughtfulness and preparation = Conscientiousness. Knowledge and market depth = Capability. Cognitive ability signals = Capability. Industry expertise strengthens fit score under Capability.

Agent architecture — browser extension vs ATS widget	RESOLVE D	All agents (Screening, Interview, Sourcing) are browser extensions that run autonomously and write results to the ATS widget. The agent does the research FOR the recruiter, not WITH the recruiter. Recruiter opens the widget and finds a complete intelligence package. The ATS widget is the display layer. The extensions are the intelligence layer. System is ATS-agnostic by design — intelligence lives in the extensions, not the ATS.
--	-----------	--

9. Agent Architecture

All three agents in the system are browser extensions that operate autonomously and write their results to the ATS. The recruiter is the decision-maker, not the researcher. The agent does the work.

This is not AI-assisted recruiting. It is AI-led research with human-in-the-loop decisions. The distinction matters — assisted implies the human is still doing the work with help. Led means the AI completes the research autonomously and returns a result for human review and decision.

Agent	Interface	When It Runs	What It Does
Screening A[i]gent	Browser extension	At application_submitted	Researches candidate autonomously across LinkedIn, employer sites, public forums, GitHub, speaking events, published articles, community mentions. Returns complete intelligence package to ATS widget: verified experience, role similarity, tech stack validation, public evidence, query flags, ask-next prompts, recommendation. Recruiter never opens a browser tab.
Interview A[i]gent	Browser extension	During interview	Runs autonomously via AI notetaker. Listens, transcribes, scores against rubric in real time. Surfaces follow-up question suggestions in ATS widget. Writes scorecard to candidate file automatically post-interview. Recalculates construct scores and generates next-steps recommendation. Recruiter focuses on the candidate — agent handles documentation.
Sourcing A[i]gent	Browser extension	At open req / search criteria trigger	Researches and surfaces candidates autonomously. Returns shortlist to recruiter with intelligence already attached — no manual research required.

The ATS widget is the display and decision layer. The browser extensions are the intelligence layer that populates it. The system is ATS-agnostic by design — the intelligence lives in the extensions, not in Greenhouse or any other ATS. Swap the ATS and the agents continue to function.

10. Stage 4 — Offer Stage & Offer A[i]gent

The offer stage is where the full signal history of the process pays off. The Offer A[i]gent draws on everything the candidate has revealed across all prior stages — compensation expectations, competing processes, priorities, constraints — and surfaces a tailored negotiation strategy before the recruiter picks up the phone.

10.1 Offer A[i]gent — What It Does

Pre-offer preparation:

- Surfaces intel from all prior conversations and signal data: compensation expectations, other interview processes in flight, location constraints, stated priorities (health insurance, equity, flexibility, growth)
- Flags visa and work authorization requirements
- Identifies relocation needs and potential assistance requirements
- Cross-references against approved budget, comp bands, approval thresholds, and current team comp makeup
- Flags role-specific timing issues — example: a Q4 sales candidate who would forfeit commissions by leaving now. Agent flags: suggest signing bonus to offset, recommend 6-8 week start delay, here is the math. Agent knows this because it has been listening across the entire process.
- Suggests negotiation tactics calibrated to what this specific candidate has signaled they care about
- Surfaces what is approved and what requires escalation before the call begins

During the offer call:

- Listens via notetaker
- Auto-populates agreed terms in real time as conversation progresses
- Flags items not yet covered that the rubric says should be addressed
- Triggers approval workflows as terms are verbally confirmed

Post-call:

- Sends for automatic approvals based on agreed and recorded terms
- Routes to DocuSign automatically — no manual handoff
- Tracks signature status and fires onboarding trigger on signed DocuSign

The recruiter walks into the offer call already knowing what this specific candidate needs to say yes — because the system has been listening and scoring across the entire process. This is not manipulation. It is preparation grounded in genuine information the candidate provided. The outcome is better for both sides: candidate gets an offer tailored to their actual situation, company closes faster with less friction.

10.2 Deliberation Phase Monitoring

Between offer_sent and offer response, the system monitors candidate engagement and nudges appropriately — or not at all, depending on the candidate's signal history.

Engagement Pattern	Signal	System Response
--------------------	--------	-----------------

High engagement — portal sessions, questions, outreach	Strong Motivation — seriously considering	No nudge needed. Let them deliberate. Surface to recruiter as positive signal.
Complete silence	Ambiguous — decided yes and waiting, or losing interest	System reads prior signal history to determine nudge vs no nudge. Candidate who has been consistently responsive = likely fine. Candidate whose signals have been declining = light recruiter touchpoint suggested.
Outreach with specific questions	Motivation + Conscientiousness — engaged and thorough	Recruiter responds personally. System surfaces any unanswered items from offer conversation.
Missing response deadline without communication	Conscientiousness negative signal	Recruiter touchpoint triggered. Deadline extension or withdrawal decision required.

10.3 Offer Stage Observations & Signals

Observation	Event	Construct	Notes
offer_response_time_hours	offer_accepted / offer_declined	Conscientiousness	How quickly candidate responded to offer. Missing deadline without communication = negative signal.
negotiation_initiated	Between offer_sent and response	Motivation	Candidate engaged with the offer rather than ignoring or instantly accepting. Engagement is positive regardless of direction.
candidate_engagement_score_deliberation	portal_session / recruiter_touchpoint	Motivation	Composite of engagement signals during deliberation window. Feeds nudge logic.
declined_reason	offer_declined	Decision Quality	Categorical: compensation / competing offer / role fit / process experience / personal / other. Feeds Decision Quality and future sourcing intelligence.
competing_offer_present	offer_declined or negotiation	Decision Quality	Boolean. If true, surfaces comp benchmarking flag and time-to-close analysis.
signing_bonus_required	Offer negotiation	Workflow	Boolean. Flags comp strategy issue — should inform future offer architecture for similar roles.
relocation_assistance_required	Offer negotiation	Workflow	Boolean + amount. Feeds total comp calculation and approval routing.
docusign_completed	offer_accepted	Workflow terminus	Triggers onboarding sequence automatically. Time from offer_sent to signed = offer cycle time for Workflow Intelligence.
offer_cycle_time_hours	offer_sent to docusign_completed	Workflow Intelligence	How long from offer to signed. Benchmark for future offer process optimization.

10.4 Updated Agent Architecture — Offer A[i]gent Added

Agent	When It Runs	What It Does Autonomously
Screening A[i]gent	application_submitted	Researches candidate autonomously. Returns complete intelligence package to ATS widget. Recruiter never opens a tab.
Interview A[i]gent	During every interview stage	Listens, transcribes, scores against rubric, surfaces follow-up suggestions, writes scorecard to ATS automatically post-interview.
Offer A[i]gent	Pre-offer through DocuSign	Surfaces prior signal intel and negotiation strategy. Auto-populates agreed terms during offer call. Routes approvals. Triggers DocuSign and onboarding on completion.
Sourcing A[i]gent	Open req / search trigger	Researches and surfaces candidates autonomously with intelligence pre-attached.

11. Stage 5 — Onboarding & The Referral Loop

Onboarding is the process terminus for candidate scoring. The DocuSign fires, the onboarding sequence triggers, and the hiring manager receives the handoff. But the relationship is not over — the 30-day check-in is where the system completes its full circle.

11.1 Onboarding Stage — Workflow Terminus

Observation	Event	Layer	Notes
docusign_completed	offer_accepted	Workflow terminus	Triggers onboarding sequence automatically. Handoff to hiring manager with 30/60/90 plan and team intro.
new_hire_survey_submitted	30 days post-start	Decision Quality	Did the candidate experience match expectations? Feeds employer brand and process quality metrics.
performance_at_30_days	30 days post-start	Decision Quality	Did behavioral signals during hiring predict actual performance? Outcome validation for scoring model.
performance_at_60_days	60 days post-start	Decision Quality	Second calibration point. Confirms or challenges 30-day read.
performance_at_90_days	90 days post-start	Decision Quality	Primary outcome signal. Did the hire work? Feeds long-term model refinement.

The 30/60/90 day performance data answers the question the Decision Quality dashboard exists to ask: did our behavioral signals actually predict success? Over time this data refines the scoring model — signals that predict performance get higher weight, signals that do not get adjusted.

11.2 The 30-Day Referral Loop — Full Circle Back to Sourcing

The 30-day check-in is not just a performance touchpoint. It is the highest-conversion moment for referral generation in the entire process. The new hire is energized, their LinkedIn is lighting up with congratulations, their network is already paying attention, and their goodwill toward the company and the recruiter is at its peak. They just experienced your candidate journey firsthand and know it was good. That is the moment to ask.

The referral itself carries signal — who someone refers tells you about their network quality, their understanding of the role, and their investment in the company's success. That is a Motivation signal that extends past the hire date.

Event	Trigger	Layer	Notes
30_day_checkin	30 days post-start	Workflow + Decision Quality	Recruiter personal touchpoint. Performance check-in and referral ask. Never automated — this is a human moment.
referral_submitted	New hire refers a contact	Sourcing + Motivation signal	Who they refer signals network quality and role understanding. Enters sourcing pipeline with referral source tagged.
referral_converted	Referred candidate advances	Decision Quality + Sourcing Intelligence	Did this hire produce quality referrals? Feeds source quality analysis and referral program ROI.

Sourcing > Application > Screen > Interview > Offer > Hire > 30-day check-in > Referral > back into Sourcing. The system is a loop, not a funnel. Every hire is a potential source. Every referral carries the signal of the person who made it.

12. Employee Intelligence & Listening A[i]gent

The candidate score does not lock at hire. It evolves into the employee lifecycle. The four constructs continue to develop as the new hire becomes an employee, with the Listening A[i]gent extending signal collection into the ongoing work environment.

The same philosophy applies post-hire as pre-hire: observable facts and behaviors only. The Listening A[i]gent does not ask how people feel. It observes what they do and say in the natural flow of work — meeting transcripts, communication patterns, sentiment trajectory — and surfaces signals without requiring self-report.

12.1 Construct Evolution Post-Hire

Construct	As Candidate	As Employee
Capability	Employer research, tech stack validation, task quality, role similarity	Performance reviews, project outcomes, skill development, promotion velocity

Conscientiousness	Response time, timeliness, no-shows, scheduling reliability, follow-through	Deadline adherence, attendance patterns, follow-through on commitments, peer feedback
Communication	Resume quality, verbal communication, message sentiment, thank you notes	360 feedback, meeting participation, written quality in work product, stakeholder management
Motivation	Portal engagement, opt-in submissions, prep investment, proactive outreach	Engagement signals, discretionary effort, internal mobility interest, voluntary learning behavior

12.2 Listening A[i]gent — What It Observes

The Listening A[i]gent monitors signals across the work environment passively — not through surveys or self-report but through the natural data generated by how people communicate, collaborate, and engage at work.

Signal Source	Construct	What It Captures
Meeting transcripts	Communication + Capability	Participation quality, question depth, clarity of contribution, leadership signals
Written communication patterns	Communication + Conscientiousness	Response time, tone, sentiment trajectory, attention to detail across correspondence
Sentiment trajectory	Motivation + Culture	Is engagement increasing or declining over time? Early attrition signal before conscious decision to leave
Culture indicators	All constructs	Is this person a culture carrier or a culture risk? Alignment between stated values and observed behavior
Discretionary effort signals	Motivation	Opt-in behavior post-hire — volunteering for stretch projects, mentoring, knowledge sharing
Peer interaction patterns	Communication + Conscientiousness	Collaboration quality, follow-through on commitments to colleagues, meeting attendance and preparation
Learning and development engagement	Motivation + Capability	Voluntary training completion, certification pursuit, skill development outside required programs

12.3 Referral Quality & Employee Sentiment

Employee sentiment at the time of the referral ask significantly affects referral quality. An employee whose sentiment trajectory is positive and stable is a high-quality referral source. An employee whose sentiment is declining is likely to refer people away from the company rather than toward it — or to broadcast the link rather than make targeted recommendations.

Referral Type	Signal Quality	Detection Method
Direct referral — specific person, targeted outreach	High	referral_type = direct. Low application volume per link (1-3). Referrer evaluated candidate personally before sending.
Broadcast referral — LinkedIn post or shared link	Low / noise	referral_type = broadcast. High application volume per link in short window. No individual vetting by referrer.

		Converts at rates closer to general inbound than true referral.
--	--	---

Detection: referral_application_count_per_link over a 48-hour window. One link producing 40+ applications is a broadcast event. One link producing 1-3 is a targeted referral. Source quality analysis must segment these — mixing broadcast and direct referrals degrades the channel quality metric and overstates referral program ROI.

12.4 Decision Quality Feedback Loop — Did the System Call It Right?

The most powerful validation in the system is the connection between pre-hire behavioral signals and post-hire performance outcomes. Over time the 30/60/90 day performance data answers the question the scoring model was always trying to answer: did what we observed during hiring predict how this person actually performs?

Question	How the System Answers It
Did Capability scores predict performance?	Compare pre-hire Capability construct scores against 90-day performance ratings. High correlation validates the signal weights. Low correlation triggers rubric review.
Did Conscientiousness scores predict retention?	Compare pre-hire Conscientiousness scores against 12-month retention. Candidates who scored low and left early = model validation. Candidates who scored high and left = dig into why.
Did Motivation scores predict engagement?	Compare pre-hire Motivation scores against post-hire engagement signals from Listening A[i]gent. High motivation candidates who disengage quickly may indicate onboarding or management failures, not hiring failures.
Did overriding the agent recommendation produce better or worse hires?	Compare performance outcomes for agent-aligned decisions vs recruiter overrides. If overrides systematically underperform, that is a Decision Quality finding requiring calibration.
Which interviewers called it right?	Compare interviewer impression scores against 90-day performance. Interviewers whose scores correlate with outcomes are reliable signal sources. Those who don't are noise — and the system can identify them without anyone having to accuse them of bias.

The system gets smarter with every hire. Construct weights that predict performance get reinforced. Signals that don't correlate with outcomes get adjusted. Interviewers whose scores are reliable get surfaced as calibration anchors. Interviewers who introduce noise get flagged for coaching — without anyone having to accuse them of bias, because the data says it.