

Listening A[i]gent

Post-Hire Talent Intelligence Layer

Design Principles, Signal Architecture, Output Layer & Triggered Actions

Diane Wilkinson | AI-Native TA Operations Portfolio | dianewilkinson.github.io | Version 4 — March 2026

1. The Core Insight

No one has more complete information about how work actually gets done than the assistant who was there when it happened.

The Listening A[i]gent observes how work unfolds in real time across every interaction — the questions asked, the ideas raised, the problems flagged, the commitments made and kept. Not self-reported. Not filtered through a manager. Observed directly, at volume, over time.

Volume is what makes this different. Any single observation is anecdotal. Hundreds of observations across months of actual work produce a signal that no manager review cycle, no annual survey, and no performance conversation has ever been able to produce: a truthful, behavioral picture of how this person actually works.

This system answers questions that were previously unanswerable. Not because the information did not exist — but because no one was in a position to observe it consistently, at scale, without the distortions that come from being a stakeholder in the outcome.

The agent has two functional layers operating as one product. The collection layer observes continuously and passively — signals accumulate from work already happening with no additional instrumentation required. The intelligence layer surfaces what was observed: to the employee as daily growth feedback, to the org as routed action. From the outside, there is only one agent. The internal architecture is not the point.

The same four constructs that predicted performance during hiring continue to develop post-hire. The scoring engine does not change. The data sources do. Pre-hire asked: what is this candidate showing us? Post-hire asks: what is this employee showing us — and what do they need to know to grow faster?

2. Daily Growth Acceleration

The primary value of the Listening A[i]gent is not periodic assessment. It is continuous, specific, behavioral feedback that helps the employee improve on a daily basis.

The employee can ask at any time: what am I showing right now? Where is my trajectory heading and what would move it? The system responds with specific, evidence-based feedback grounded in observed behavior — not manager opinion, not self-report, not once-a-year reflections after the moment has long passed.

The agent was there when the work happened. It does not need the employee to describe what they did. The conversation starts from evidence, not narrative. Growth accelerates because the feedback loop collapses from annual to continuous.

Instead of a periodic judgment handed down from above, the employee has a real-time, always-available read on their own behavioral profile — and a direct path to improving it today.

3. The Agency Problem — Why Idea Routing Matters

Business schools have a name for the structural failure this system is designed to solve: the agency problem. Principals — shareholders, executives, the organization — want value created. Agents — managers, employees — have their own incentives that do not always align with value creation.

In practice: the person closest to the problem — doing the actual work, seeing the actual inefficiencies, generating the actual insight — cannot get that insight to the person with authority to act on it. The manager in the middle may not have the technical depth to evaluate the idea, the organizational standing to champion it, or the personal incentive to elevate someone whose capabilities might outgrow the role they need filled.

The result is systemic signal loss. Ideas that would create genuine competitive advantage die in the space between the person who had them and the person who could have used them. Not because of bad intent. Because of structural misalignment.

No hack days required. No suggestion boxes. No quarterly innovation initiatives. The ideas are already being generated by the people closest to the work. The system routes them directly to where they can be evaluated on merit — bypassing the bottleneck structurally, not through culture change.

An employee raises an idea. The system evaluates whether it is genuinely novel, identifies who has context and authority to evaluate it, and routes it directly with full context attached. Outcome

tracking closes the loop — whether an idea was picked up, modified, deprioritized, or dropped is logged and feeds Decision Quality. Over time: is innovation routing in this org merit-based, or does it follow the existing political topology?

4. Equity By Design

Not every employee is an ideas generator. Not every role calls for visible meeting participation. Not every valuable contributor is the loudest voice in the room. A system that scores on contribution visibility rather than contribution quality reproduces the same bias it is designed to eliminate.

- Absence of meeting participation is neutral, not negative. The system does not penalize introverts, non-native speakers, or people without positional power.
- Quality of contribution when contribution occurs is what is measured — not frequency.
- Credit follows the originating signal. If an idea was raised by one person and adopted by another, the routing layer has the timestamp. Attribution is traceable.
- Role-appropriate scoring means construct weights are configured to what the role actually requires. An operations coordinator is not penalized for low idea submission frequency.
- Not every org is made of innovators. People who maintain, execute, and stabilize are equally essential. The scoring configuration must reflect that.

The system advocates for the employee who does excellent work quietly. That is who it is built for.

5. Performance Reviews — Reframed

The manager's job in performance management is not to generate the assessment. The Listening A[i]gent already did that. The manager's job is to be the human in the loop — to validate what the data shows, to have the forward-looking conversation, to confirm the plan, to encourage the trajectory.

That is a fundamentally different role than scoring someone on a 1-5 scale based on recency bias and political relationship. The Listening A[i]gent removes the part of performance management that managers were never well-positioned to do — objective behavioral assessment — and replaces it with something they are well-positioned to do: validate growth, cheer people on to greater things, and commit to a forward plan.

The review conversation is not a judgment. It is a calibration. Both parties arrive with the same evidence base. The manager's role is to confirm, encourage, and plan — not to produce a score from memory.

5.1 The Manager Conversation Tool

The manager has a dedicated guided workflow surface within the Talent Intelligence dashboard that fires at review cadence. It is not a scoring form. It is a structured conversation guide that:

- Presents the signal picture the Listening A[i]gent assembled — construct scores, trajectory, acceleration, key signal highlights — so the manager arrives prepared
- Prompts through human confirmation points: does this match what you observed? What do you want to reinforce? What forward action are you committing to?
- Logs manager confirmations and forward commitments as human-in-the-loop decisions that feed Decision Quality
- Surfaces growth conversation prompts calibrated to the employee's specific signal profile and trajectory
- Records the agreed plan — development focus, stretch opportunities, timeline — as a structured output that both parties can reference

The manager becomes a validator and a coach rather than a judge. The tool makes that role change structural rather than cultural — you do not have to ask managers to be better, you change what the tool asks them to do. The managers who resist that shift are the ones whose value was concentrated in the gatekeeping function. The ones who thrive are the ones who were always good at the human side.

6. Signal Architecture — Post-Hire

6.1 How Signals Fire

The same three-layer model applies post-hire: Event, Observation, Signal. A single event can generate multiple observations, each normalizing to a separate signal column. There is no double-counting — each signal measures a distinct dimension of the same moment.

Example Event	Observations Generated	Signals Fired
Meeting — deep technical question asked	knowledge_depth_raw, question_quality_raw	knowledge_depth_signal (Capability) + question_quality_signal (Capability)
Same meeting — how the question was phrased	verbal_communication_raw	question_delivery_signal (Communication)
Written deliverable submitted	work_product_quality_raw, written_quality_raw	work_product_quality_signal (Capability) + written_quality_signal (Communication)
Colleague interaction — fast clear response	response_time_hours, message_sentiment_raw	responsiveness_signal (Communication) + message_sentiment_signal (Communication)
Stretch assignment completed voluntarily	task_quality_raw, opt_in_boolean	stretch_project_signal (Motivation) + work_product_quality_signal (Capability)

Meeting event taxonomy — standup vs strategy session vs 1:1 vs cross-functional — is an open item. See Section 11.

6.2 Capability Signals (40% default — configurable per role)

Signal	Post-Hire Source
work_product_quality_signal	Quality of actual deliverables — writing, analysis, code, outputs
knowledge_depth_signal	Confidence and depth of domain knowledge in discussion — how far below the surface
question_quality_signal	Insight and depth of questions asked — substance, not delivery
idea_originality_signal	Novelty of ideas surfaced through routing layer vs what is already in flight
peer_acknowledgment_signal	Ideas or contributions built on by colleagues — credit follows the originating signal
skill_application_signal	Observable use of claimed skills in actual work context
role_expansion_signal	Taking on adjacent scope beyond defined role — opt-in signal, not required

6.3 Conscientiousness Signals (25% default — configurable per role)

Signal	Post-Hire Source
deadline_adherence_signal	Task completion relative to committed deadlines
commitment_follow_through_signal	Commitments made in meetings — completed vs dropped, traceable from transcript
meeting_reliability_signal	Attendance and punctuality — employee-initiated absences only
meeting_preparation_signal	Evidence of preparation in participation quality — not attendance alone
response_delay_signal	Response time to colleague communications — carries forward from pre-hire
work_product_completeness_signal	Thoroughness and follow-through in deliverable quality

6.4 Motivation Signals (20% default — configurable per role)

Signal	Post-Hire Source
voluntary_learning_signal	Opt-in training, certification pursuit, learning beyond required programs

Signal	Post-Hire Source
stretch_project_signal	Volunteering for stretch assignments — boolean per type, not additive volume
knowledge_sharing_signal	Proactive contribution to team knowledge — not required, self-directed
initiative_signal	Self-directed work beyond assigned scope — carries forward from pre-hire logic
idea_submission_signal	Frequency and consistency of ideas submitted — configurable weight per role
proactive_communication_signal	Unprompted updates, early flagging of blockers
sentiment_trajectory_signal	Engagement trend over rolling window — rising or declining

Floor of 0.2 applied if employment_status = active_employee and tenure > 0 days.

6.5 Communication Signals (15% default — configurable per role)

Signal	Post-Hire Source
responsiveness_signal	Response time to colleagues — carries forward from pre-hire
message_sentiment_signal	Sentiment of written communication across channels
written_quality_signal	Quality of written work product — grammar, clarity, precision
question_delivery_signal	Clarity and fluency of how questions are phrased — delivery, not substance
verbal_communication_signal	Meeting participation quality — tone, fluency, interpersonal effectiveness

question_quality_signal (substance) lives in Capability. question_delivery_signal (delivery) lives in Communication. Same event, two observations, two construct contributions. Not double-counting.

7. Acceleration Modifier

Acceleration measures the rate of score improvement over time — not where the employee is today, but how fast they are getting there. This is the behavioral observable of general mental ability (GMA), the single strongest predictor of job performance across all roles (Schmidt & Hunter, 1998).

Acceleration is a modifier alongside the four construct scores. It does not feed into the overall weighted score. It informs decisions the overall score cannot make alone — who is performing now versus who will be ready for more scope in 12 months.

Pattern	Acceleration Read	Decision Implication
High score, rising fast	Strong momentum	High performer on trajectory — retention priority
High score, flat	Established performer	Possibly plateauing — stretch assignment signal
Low score, rising fast	High potential	Early tenure or onboarding lag — invest and watch
Low score, flat or falling	Concern signal	Performance conversation warranted

Computation uses `score_history.csv`, which already contains cumulative scores at every event timestamp. The slope of the overall score curve is the acceleration signal. Infrastructure exists — requires a new cell in `Hex` and a column in `candidate_profile`.

8. Output Layer — Who Sees What

Each audience receives a different view of the same underlying signal data. The employee's individual behavioral signals are never exposed to their manager in raw form — managers see patterns and flags, not individual score breakdowns.

Audience	What They See	Where	Governance
Employee	Own construct scores, trajectory, acceleration, specific behavioral feedback, gap identification, L&D suggestions	Chat interface — always available	Employee-controlled. Own data only.
Manager	Team signal patterns, construct gaps, individual trajectories, succession readiness flags, attrition risk alerts. Manager conversation tool at review cadence.	Talent Intelligence dashboard — team view + manager conversation tool	Patterns and flags only. Not individual raw scores by default.
HR / People Ops	Aggregate workforce intelligence, team health by construct, L&D routing queue, attrition risk by department, headcount gap signals	Workflow Intelligence + Decision Quality	Aggregate at team level. Individual access requires HR role permission.
Executive / Leadership	Org-level Decision Quality: manager reliability, idea routing merit, succession pipeline depth, performance review calibration patterns	Decision Quality dashboard	Org-level patterns only. No individual employee data.

Audience	What They See	Where	Governance
L&D	Individual gap-to-program routing queue — construct area, gap signal, recommended program, outcome tracking. Build-vs-buy flag when no program matches.	Workflow Intelligence — L&D routing queue	Gap type and routing only. No underlying scores.

9. Dashboard Routing

9.1 Talent Intelligence Dashboard

Extends the pre-hire candidate view into the employee lifecycle. The same construct scores and acceleration modifier continue as the employee's longitudinal behavioral profile.

- Employee view: own scores, trajectory, percentile rank within team and department, growth recommendations
- Manager team view: team construct profiles, gap patterns, succession readiness indicators, attrition risk flags
- Manager conversation tool: guided review workflow that fires at cadence — presents signal picture, prompts human confirmations, logs forward commitments
- Team comparison: how does this employee's behavioral profile compare to existing team members at same tenure and role level?
- Pre-hire to post-hire continuity: did pre-hire signal scores predict post-hire performance? This is the validation layer that makes the scoring model self-improving.

9.2 Decision Quality Dashboard

- Performance review calibration: do manager confirmations align with behavioral signal scores? Divergence is a finding.
- Succession pipeline: acceleration patterns across the org — who is ready for more scope, who is plateauing, what is the pipeline depth by role and department
- Idea routing merit analysis: are ideas that get acted on coming from employees with strongest signal profiles, or does routing follow positional topology?
- Manager reliability: do managers whose teams show strong signal profiles receive more autonomy? Do managers whose forward commitments consistently go unmet get flagged?
- Pre-hire signal validation: construct scores at hire correlated with performance_rating at 30/60/90 days

9.3 Workflow Intelligence Dashboard

- Attrition risk queue: employees with declining sentiment trajectory — surfaced before they leave
- Backfill triggers: attrition risk flag + role criticality + pipeline depth = backfill decision surface
- Headcount gap signals: construct gaps across a team pointing to a missing capability — surfaces before a req is formally opened
- L&D routing queue: volume of identified gaps by type, department, and recommended program. Build-vs-buy flag when no existing program matches a gap type.
- Succession-driven role planning: acceleration patterns suggesting someone is ready for a senior role that does not yet exist

10. Triggered Actions

Seven categories of automated action fire from post-hire signal data.

10.1 L&D Routing

Field	Detail
Fires when	Signal gap identified: construct score significantly below role-appropriate threshold, sustained over rolling window
What it does	Identifies gap by construct and signal. Surfaces targeted L&D program matched to gap type. Routes to L&D queue with context: employee tenure, role, gap signal, recommended program.
Who is notified	Employee via chat (here is what the data shows and here is a path). L&D team (routing queue). Manager (flag only).
No matching program	Gap is flagged as a build-vs-buy signal in Workflow Intelligence. L&D or HR determines whether to source an external program, build internally, or address through role redesign.
Outcome tracked	Did the employee engage with the program? Did the signal gap close? Fed back as L&D program effectiveness data.

10.2 Attrition Risk Alert

Field	Detail
Fires when	sentiment_trajectory_signal declining over 60-day rolling window + tenure pattern consistent with pre-departure behavior
What it does	Surfaces attrition risk flag to HR. Generates recommended retention action: compensation review, role scope conversation, manager calibration.
Who is notified	HR / People Ops (risk flag + recommended action). Manager (flag only). Employee is not notified — this is a retention intelligence signal, not disciplinary.

Field	Detail
Outcome tracked	Was retention action taken? Did sentiment recover? If employee departs, outcome logged for pattern validation.

10.3 Succession Flag

Field	Detail
Fires when	High acceleration modifier + high overall score + role_expansion_signal present + tenure within succession-eligible window
What it does	Surfaces succession-ready flag. Identifies construct profile and role level the employee is signaling readiness for. Cross-references open roles and headcount plans. Feeds succession pipeline view in Decision Quality.
Who is notified	HR / People Ops. Executive / Leadership (succession pipeline view). Manager (flag — gives them the conversation to have).
Outcome tracked	Was the employee considered for the next opportunity? If not acted on and employee subsequently departs, that is a Decision Quality finding.

10.4 Performance Review Preparation

Field	Detail
Fires when	Performance review cycle trigger — configurable cadence per org
What it does	Surfaces behavioral signal summary to both employee and manager before the review conversation. Activates manager conversation tool in Talent Intelligence dashboard.
Who is notified	Employee (chat: here is what your behavioral signals show going into this review). Manager (dashboard: signal summary with trajectory, acceleration, and conversation guide).
Outcome tracked	Manager confirmations and forward commitments logged to Decision Quality. Divergence between signal data and manager observations is a finding — not a correction, a conversation.

10.5 Idea Routing

Field	Detail
Fires when	Employee submits idea through chat interface
What it does	Evaluates novelty vs in-flight work. Identifies right stakeholder by domain and authority. Routes with full context. Communicates next steps and expected response timeline back to employee. If idea overlaps with in-flight work, connects employee to the relevant team.
Who is notified	Routing target (idea + context). Employee (acknowledgment + timeline).

Field	Detail
Outcome tracked	Was the idea acted on? By whom? How long did routing take? Feeds Decision Quality innovation routing analysis.

10.6 Headcount Planning Signal

Field	Detail
Fires when	Construct gap pattern sustained across a team (multiple employees showing the same gap) OR attrition risk flag on a role with no internal succession candidate OR acceleration pattern suggesting a role level does not yet exist for a ready employee
What it does	Surfaces headcount gap signal to HR and leadership with context: gap type, affected team, urgency level, whether gap can be addressed through L&D or requires a hire.
Who is notified	HR / People Ops (headcount signal + recommended action). Executive / Leadership if gap is at senior level or cross-functional.
Outcome tracked	Was a req opened? Was a learning intervention attempted first? Did it close the gap? Outcome logged to Workflow Intelligence for future calibration.

10.7 Compensation Equity Flag

Field	Detail
Fires when	Promotion or compensation review where candidate scores lower than existing team members in same role/level but market rate requires offering above current team compensation
What it does	Surfaces equity risk flag before the decision is made. Not a block — a flag. The decision is made with full information.
Who is notified	HR / People Ops. Decision owner. Carries forward from pre-hire offer stage logic.
Outcome tracked	Was the flag acknowledged? What was the decision? If inequity was accepted, logged to Decision Quality for future team health monitoring.

11. Learning A[i]gent — On the Horizon

The Listening A[i]gent observes and surfaces behavioral performance intelligence. The Learning A[i]gent is a distinct future agent that does something architecturally different: it analyzes skill gaps at individual and org level, models the reskill vs upskill vs hire decision, and feeds workforce planning with predictive hiring needs before a gap becomes a crisis.

The distinction matters because the Listening A[i]gent answers what is happening now. The Learning A[i]gent answers what the org will need in 12-18 months and whether it can develop

that capability internally or must hire for it. That is a forward-looking, analytical function that draws on Listening A[i]gent outputs as inputs but operates on a different question entirely.

Function	Listening A[i]gent	Learning A[i]gent
Timeframe	Present — what is happening now	Forward — what will be needed
Question answered	What is this employee showing us?	Can we develop this capability internally or must we hire?
Primary output	Behavioral signal scores, growth feedback, idea routing	Skill gap analysis, reskill/upskill/hire recommendation, predictive req pipeline
Feeds into	Talent Intelligence, Decision Quality, Workflow Intelligence	Workforce planning, L&D program design, sourcing strategy
Status	Designed — partially demonstrable now	Future agent — specced, not yet built

The Learning A[i]gent is documented here so the architecture does not have to be retrofitted when it gets built. The Listening A[i]gent's L&D routing function — gap identification and program matching — is the seed of what the Learning A[i]gent will eventually formalize at scale.

12. AI Use Is Not a Signal Problem

If an employee uses AI to produce higher quality work, that output still represents genuine signal. The system does not penalize AI assistance. Quality is quality regardless of how it was produced.

The behavioral signal that matters is whether the employee understands, can build on, and can defend the work. That surfaces through question quality in discussion, response to challenge, idea development over time. An employee who ships AI output without genuine engagement will show a divergence between output quality signals and discussion quality signals across hundreds of observations.

AI can optimize a single output. It cannot sustain a false signal pattern across months of behavioral data without the underlying capability to support it.

13. Open Items

Item	Decision Needed
Meeting event taxonomy	Standup vs strategy session vs 1:1 vs cross-functional — different signal profiles per type or context inference from transcript?

Item	Decision Needed
idea_submission_signal weight by role	Roles where idea generation is not a primary function — floor weight vs zero?
sentiment_trajectory_signal sourcing	Enterprise data sources priority order: Slack, email, meeting transcripts — coverage requirements.
Credit attribution for adopted ideas	Technical implementation for tracing originating signal when an idea from employee A is built on by employee B.
Acceleration rolling window	30 days, 90 days, or full tenure? Different windows answer different questions.
Manager view access escalation	Default: patterns and flags only. What access path exists when investigation is warranted?
Attrition risk threshold calibration	Slope and duration on sentiment_trajectory_signal that constitutes a risk flag — requires historical data.
Manager conversation tool cadence	Annual, semi-annual, or continuous with formal trigger points? Org-configurable or system default?

14. Built vs. Conceptual

Component	Status	Tool
Four-construct scoring engine	Built	Google Sheets + Apps Script
Pre-hire signal collection	Built	signals.csv + Hex
score_history cumulative logic	Built	Hex scoring engine cell
Acceleration modifier computation	Ready to build	Hex new cell + candidate_profile column
Post-hire performance correlation	Designed — synthetic data	Hex + employment_status + performance_rating
Opt-in assessment chat interface	Demonstrable now	Claude API artifact
Idea routing demo	Demonstrable now	Airtable + Claude API
Manager conversation tool	Conceptual	Guided workflow within Talent Intelligence dashboard
L&D routing trigger	Conceptual	Requires signal gap threshold calibration
Attrition risk alert	Conceptual	Requires historical departure pattern data
Succession flag	Conceptual	Acceleration modifier build required first
Headcount planning signal	Conceptual	Requires team-level gap pattern detection
Performance review prep trigger	Conceptual	Straightforward once signal data is live

Component	Status	Tool
Compensation equity flag	Designed	Carries forward from pre-hire offer stage logic
Listening A[i]gent signal collection	Conceptual	Requires enterprise integration — Slack, transcript pipeline
Innovation pattern detection	Conceptual	Designed, not yet built
Decision Quality feedback loop	Designed	Hex — pre-hire score vs post-hire performance_rating
Learning A[i]gent	Future agent	Specced — not yet built

15. What This System Is Not

- Not surveillance — signals are collected from work already happening, not additional monitoring instrumentation
- Not a replacement for manager judgment — it gives managers better evidence and a better tool for the human work they are actually good at
- Not anti-AI — AI-assisted work that produces genuine quality is the goal, not the problem
- Not opinion-based — manager judgment is excluded from construct scores for the same reason it is excluded pre-hire
- Not biased toward the visible — absence of outspoken meeting participation is neutral, not penalized
- Not a one-size scoring model — construct weights are configurable per role, per team, per org
- Not a suggestion box — the routing layer closes the loop; outcome tracking is built in
- Not a new layer of pressure on people who maintain and execute — not every role is an innovation role
- Not a performance review replacement — it is a performance review upgrade. The review still happens. The manager still matters. The evidence base is now honest.

The behavioral signal score tells you what the employee showed. The manager conversation tells you what comes next. Keeping assessment and forward action as distinct functions is what makes both better.

v4 — March 2026. Product renamed Listening A[i]gent throughout. Listening and Insights merged into one product — collection layer is infrastructure, not a separate named agent. Performance reviews reframed: manager validates, coaches, and commits to forward action rather than scoring. Manager conversation tool added as guided workflow surface within Talent Intelligence dashboard. L&D gap with no matching program documented as build-vs-buy signal. Headcount planning promoted to triggered action (10.6). Succession pipeline view described as Decision Quality output. Learning A[i]gent added as named future agent (Section 11). Open items expanded to eight.